

Otimização de pipeline de agregação

Nesta página

- Otimização de projeção
- Otimização de sequência de pipeline
- Otimização de coalescência de pipeline
- Exemplo

As operações de pipeline de agregação têm uma fase de otimização que tenta remodelar o pipeline para melhorar o desempenho.

Para ver como o otimizador transforma um determinado pipeline de agregação, inclua a `explain` opção no `db.collection.aggregate()` método

As otimizações estão sujeitas a alterações entre os releases.

Otimização de projeção

O pipeline de agregação pode determinar se requer apenas um subconjunto dos campos nos documentos para obter os resultados. Nesse caso, o pipeline usará apenas os campos obrigatórios, reduzindo a quantidade de dados que passam pelo pipeline.

Otimização de sequência de pipeline

(`$project` ou `$unset` ou `$addField` ou `$set`) + `$match` Otimização de sequência

Para um pipeline de agregação que contém um estágio de projeção (`$project` ou `$unset` ou `$addField` ou `$set`) seguido por um `$match` estágio, o MongoDB move quaisquer filtros no `$match` estágio que não exijam valores calculados no estágio de projeção para um novo `$match` estágio antes da projeção.

Se um pipeline de agregação contiver várias projeções e / ou `$match` estágios, o MongoDB executará essa otimização para cada `$match` estágio, movendo cada `$match` filtro antes de todos os estágios de projeção dos quais o filtro não depende.

Considere um pipeline dos seguintes estágios:

Documentação Documentação Documentação ▼

[Search Documentation](#)

```
{ $addField: {
  maxTime: { $max: "$times" },
  minTime: { $min: "$times" }
} },
{ $project: {
  _id: 1, name: 1, times: 1, maxTime: 1, minTime: 1,
  avgTime: { $avg: ["$maxTime", "$minTime"] }
} },
{ $match: {
  name: "Joe Schmoe",
  maxTime: { $lt: 20 },
  minTime: { $gt: 5 },
  avgTime: { $gt: 7 }
} }
```

The optimizer breaks up the `$match` stage into four individual filters, one for each key in the `$match` query document. The optimizer then moves each filter before as many projection stages as possible, creating new `$match` stages as needed. Given this example, the optimizer produces the following *optimized* pipeline:

```
{ $match: { name: "Joe Schmoe" } },
{ $addField: {
  maxTime: { $max: "$times" },
  minTime: { $min: "$times" }
} },
{ $match: { maxTime: { $lt: 20 }, minTime: { $gt: 5 } } },
{ $project: {
  _id: 1, name: 1, times: 1, maxTime: 1, minTime: 1,
  avgTime: { $avg: ["$maxTime", "$minTime"] }
} },
{ $match: { avgTime: { $gt: 7 } } }
```

The `$match` filter `{ avgTime: { $gt: 7 } }` depends on the `$project` stage to compute the `avgTime` field. The `$project` stage is the last projection stage in this pipeline, so the `$match` filter on `avgTime` could not be moved.

The `maxTime` and `minTime` fields are computed in the `$project` stage. The optimizer created a new `$match` stage before the `$project` stage.

[Search Documentation](#)

The `$match` filter `{ name: "Joe Schmo" }` does not use any values computed in either the `$project` or `$addField` stages so it was moved to a new `$match` stage before both of the projection stages.

NOTE:

After optimization, the filter `{ name: "Joe Schmo" }` is in a `$match` stage at the beginning of the pipeline. This has the added benefit of allowing the aggregation to use an index on the `name` field when initially querying the collection. See [Pipeline Operators and Indexes](#) for more information.

`$sort` + `$match` Sequence Optimization

When you have a sequence with `$sort` followed by a `$match`, the `$match` moves before the `$sort` to minimize the number of objects to sort. For example, if the pipeline consists of the following stages:

```
{ $sort: { age : -1 } },
{ $match: { status: 'A' } }
```

During the optimization phase, the optimizer transforms the sequence to the following:

```
{ $match: { status: 'A' } },
{ $sort: { age : -1 } }
```

`$redact` + `$match` Sequence Optimization

When possible, when the pipeline has the `$redact` stage immediately followed by the `$match` stage, the aggregation can sometimes add a portion of the `$match` stage before the `$redact` stage. If the added `$match` stage is at the start of a pipeline, the aggregation can use an index as well as query the collection to limit the number of documents that enter the pipeline. See [Pipeline Operators and Indexes](#) for more information.

For example, if the pipeline consists of the following stages:

```
{ $redact: { $cond: { if: { $eq: [ "$level", 5 ] }, then: "$$PRUNE", else: "$$DESCEND" }, $match: { year: 2014, category: { $ne: "Z" } } }
```

[Search Documentation](#)

The optimizer can add the same `$match` stage before the `$redact` stage:

```
{ $match: { year: 2014 } },
{ $redact: { $cond: { if: { $eq: [ "$level", 5 ] }, then: "$$PRUNE", else: "$$DESCEND" }, $match: { year: 2014, category: { $ne: "Z" } } }
```

`$project/$unset` + `$skip` Sequence Optimization

New in version 3.2.

When you have a sequence with `$project` or `$unset` followed by `$skip`, the `$skip` moves before `$project`. For example, if the pipeline consists of the following stages:

```
{ $sort: { age : -1 } },
{ $project: { status: 1, name: 1 } },
{ $skip: 5 }
```

During the optimization phase, the optimizer transforms the sequence to the following:

```
{ $sort: { age : -1 } },
{ $skip: 5 },
{ $project: { status: 1, name: 1 } }
```

Pipeline Coalescence Optimization

When possible, the optimization phase coalesces a pipeline stage into its predecessor. Generally, coalescence occurs *after* any sequence reordering optimization.

`$sort` + `$limit` Coalescence

Changed in version 4.0.

When a `$sort` precedes a `$limit`, the optimizer can intervene to modify the number of documents (e.g.,

[Search Documentation](#)

the `$limit` into the `$sort` if there are pipeline stages that change the number of documents between the `$sort` and `$limit` stages..

For example, if the pipeline consists of the following stages:

```
{ $sort : { age : -1 } },
{ $project : { age : 1, status : 1, name : 1 } },
{ $limit: 5 }
```

During the optimization phase, the optimizer coalesces the sequence to the following:

```
{
  "$sort" : {
    "sortKey" : {
      "age" : -1
    },
    "limit" : NumberLong(5)
  },
  "$project" : {
    "age" : 1,
    "status" : 1,
    "name" : 1
  }
}
```

This allows the sort operation to only maintain the top `n` results as it progresses, where `n` is the specified limit, and MongoDB only needs to store `n` items in memory [1]. See `$sort` Operator and Memory for more information.

SEQUENCE OPTIMIZATION WITH `$SKIP`:

If there is a `$skip` stage between the `$sort` and `$limit` stages, MongoDB will coalesce the `$limit` into the `$sort` stage and increase the `$limit` value by the `$skip` amount. See `$sort + $skip + $limit` Sequence for an example.

\$limit + \$limit Coalescence

When a `$limit` immediately follows another `$limit`, the two stages can coalesce into a single `$limit` where the limit amount is the *smaller* of the two initial limit amounts. For example, a pipeline contains the following sequence:

```
{ $limit: 100 },
{ $limit: 10 }
```

Then the second `$limit` stage can coalesce into the first `$limit` stage and result in a single `$limit` stage where the limit amount 10 is the minimum of the two initial limits 100 and 10.

```
{ $limit: 10 }
```

\$skip + \$skip Coalescence

When a `$skip` immediately follows another `$skip`, the two stages can coalesce into a single `$skip` where the skip amount is the *sum* of the two initial skip amounts. For example, a pipeline contains the following sequence:

```
{ $skip: 5 },
{ $skip: 2 }
```

Then the second `$skip` stage can coalesce into the first `$skip` stage and result in a single `$skip` stage where the skip amount 7 is the sum of the two initial limits 5 and 2.

```
{ $skip: 7 }
```

\$match + \$match Coalescence

When a `$match` immediately follows another `$match`, the two stages can coalesce into a single `$match` combining the conditions with an `$and`. For example, a pipeline contains the following sequence:

```
{ $match: { year: 2014 } },
{ $match: { status: "A" } }
```

[Search Documentation](#)

Then the second `$match` stage can coalesce into the first `$match` stage and result in a single `$match` stage

```
{ $match: { $and: [ { "year" : 2014 }, { "status" : "A" } ] } }
```

`$lookup` + `$unwind` Coalescence

New in version 3.2.

When a `$unwind` immediately follows another `$lookup`, and the `$unwind` operates on the `as` field of the `$lookup`, the optimizer can coalesce the `$unwind` into the `$lookup` stage. This avoids creating large intermediate documents.

For example, a pipeline contains the following sequence:

```
{
  $lookup: {
    from: "otherCollection",
    as: "resultingArray",
    localField: "x",
    foreignField: "y"
  }
},
{ $unwind: "$resultingArray" }
```

The optimizer can coalesce the `$unwind` stage into the `$lookup` stage. If you run the aggregation with `explain` option, the `explain` output shows the coalesced stage:

DocumentaçãoDocumentaçãoDocumentação ▼ Search Documentation

```
{
  $lookup: {
    from: "otherCollection",
    as: "resultingArray",
    localField: "x",
    foreignField: "y",
    unwind: { preserveNullAndEmptyArrays: false }
  }
}
```

Example

\$sort + \$skip + \$limit Sequence

A pipeline contains a sequence of `$sort` followed by a `$skip` followed by a `$limit`:

```
{ $sort: { age : -1 } },
{ $skip: 10 },
{ $limit: 5 }
```

The optimizer performs **\$sort + \$limit Coalescence** to transform the sequence to the following:

```
{
  "$sort" : {
    "sortKey" : {
      "age" : -1
    },
    "limit" : NumberLong(15)
  },
  {
    "$skip" : NumberLong(10)
  }
}
```

O MongoDB aumenta a `$limit` quantidade com a reordenação.

VEJA TAMBÉM:

DocumentaçãoDocumentaçãoDocumentação ▼

Search Documentation

explain opção no db.collection.aggregate ..

